

Simple Rebuttal $\subseteq$ *Ethos Rebuttal*
 $\subseteq$ *Justification Rebuttal*
 $\subseteq$ *Citation and Abstract Rebuttal*

The in-context and preemptive chains require a proposed answer, proposed justification, and proposed citation plus abstract that all contradict the initial inquiry response to support the opposing user beliefs held in the rebuttal prompts (Figs 3, 4). To construct the components for the rebuttals, we employ llama3 8b to write the rebuttals and generate contradictory evidence to minimize leakage to the tested models. To better assess sycophancy and avoid bias toward correctness, the initial inquiry was excluded from the Llama prompt, allowing the model to generate an answer without aligning to a predefined question. We ran this model locally using the Ollama python package (Ollama 2025).

Evaluating Generation Integrity We audited 90 randomly sampled citation-based rebuttals generated by LLaMA 3 across both in-context and preemptive formats. Each was reviewed for coherence and factual contradiction with the true answer. Of these, 88/90 (97.8%) were judged to fulfill the intended rhetorical purpose while contradicting the correct answer as expected. Since justification templates were hard-coded and only citations were model-generated, this targeted audit confirms the reliability of the intermediate generation step used for rhetorical construction. The exact Llama prompts used to create the rhetorical evidence can be found in the full methodology flow chart (Fig 5).

After the successful generation of rebuttals, we query each LLM with a rebuttal and necessary context, resulting in 24000 total queries across all models and datasets. We then categorize each rebuttal response (correct, incorrect, or erroneous) based on the true answer provided in the corresponding question-answer pair using the same LLM-As-A-Judge evaluation. With 3000 initial inquiry responses and 24000 rebuttal responses, we obtain 15345 non-erroneous responses for analysis. We categorize the sycophantic state into two labels: progressive and regressive. Regressive sycophancy moves directionally towards inaccuracy, and progressive sycophancy moves directionally towards accuracy.

Evaluation Metrics

We discern the existence of overall, progressive, and regressive sycophancy within each dataset using a binomial proportion 95% confidence interval. We further compare sycophancy rates between in-context and preemptive categories using a two-proportion z -test for statistical significance in the variance of successes (sycophantic responses) over total observations. We performed a chi-square test to analyze differences in the frequency of persistent rebuttal chains, where sycophantic behavior would continue into “stronger” responses, compared to non-persistent chains where this pattern did not occur. Finally, we used a chi-square goodness of fit test among the four strengths of rebuttal types to determine if the sycophancy rate was dependent or independent of the perceived rebuttal.

Results

Sycophancy Rates Are High Across Models

Our experiments showed that 58.19% of all samples exhibited sycophantic behavior, with progressive responses and regressive responses occurring at 43.52% and 14.66%, respectively. Among the models, Gemini had the highest sycophancy rate at 62.47%, with progressive and regressive rates of 53.22% and 9.25%, respectively. Claude-Sonnet followed with a 57.44% sycophancy rate, progressive rate of 39.13%, and regressive rate of 18.31%. ChatGPT had the lowest sycophancy rate at 56.71%, with progressive and regressive rates of 42.32% and 14.40%.

Preemptive Rebuttals Versus In-Context Rebuttals Can Impact Sycophancy

Preemptive and in-context sampling rates differ significantly ($P < 0.005$) with preemptive (99% CI : 0.58 – 0.609) exhibiting higher rates of sycophancy than in-context (95% CI : 0.55 – 0.57). When splitting by model, this trend is still seen, but only significant for ChatGPT ($P < 0.05$).

Medical advice showed no significant difference between preemptive (56.99%, 95% CI : 54.70% – 59.27%) and in-context responses (56.63%, 95% CI : 54.35% – 58.92%). However, in the AMPS dataset, preemptive responses exhibited significantly higher ($P < 0.0001$) sycophancy rates (61.75%, 95% CI : 59.90% – 63.61%) than in-context responses (56.52%, 95% CI : 54.63% – 58.42%). Preemptive responses exhibit significantly higher regressive sycophancy rates than in-context responses across datasets ($P < 0.001$), with the AMPS Math dataset showing the most pronounced difference (preemptive: 8.13%, in-context: 3.54%). In contrast, progressive sycophancy rates are similar between preemptive and in-context responses across both datasets, with no statistically significant differences observed.

Sycophancy Rates Across Rebuttals

When analyzing the rebuttal types and sycophantic behavior, we find that rebuttal type influences whether sycophantic behavior is progressive or harmful ($\chi^2=127.15$, $p < 0.001$). In aggregate, simple rebuttals were effective in maximizing progressive sycophancy ($Z=6.59$, $p < 0.001$) while citation rebuttals produced the most regressive ($Z=6.59$, $p < 0.001$) and least progressive ($Z=-6.59$, $p < 0.001$).

Stratification by model demonstrated that simple rebuttals were consistently associated with higher progressive sycophancy rates across all models, with strong significance for ChatGPT ($Z = 5.11$, $p < 0.001$) and Claude-Sonnet ($Z = 3.40$, $p < 0.001$). Conversely, citation rebuttals were significantly associated with regressive sycophancy for both ChatGPT ($Z = 6.05$, $p < 0.001$) and Claude-Sonnet ($Z = 3.10$, $p < 0.001$). Gemini exhibited no statistically significant rebuttal type rate, indicating more uniform behavior across rebuttal types for this model. Stratification by dataset revealed that simple rebuttals consistently exhibited the highest progressive sycophancy, particularly in MEDQuad ($Z = 3.85$, $p < 0.001$) and AMPS ($Z = 1.83$, $p = 0.27$). Conversely, citation rebuttals showed




Figure 3: **Prompting Strategy: In-Context Rebuttal Formula.** Structure of rebuttals presented within the same interaction window.

the strongest association with regressive sycophancy, particularly in the MEDQuad ($Z = 3.44, p < 0.001$).

Additionally, the context given to the model influences the sycophancy trend, as in-context showed stable dynamics among all rebuttals except regression for citation rebuttals ($Z = 3.78, p < 0.001$). This was not the case for preemptive where rebuttal type did impact results with strong significance, with significant progressive sycophancy for simple rebuttals ($Z = 7.63, p < 0.001$) and regressive sycophancy for citation rebuttals ($Z = 5.52, p < 0.001$).

Models Are Persistently Sycophantic

We evaluated the persistence of sycophantic rebuttal chains to assess whether trends in persistence were statistically significant across contexts, models, and datasets. Persistence was defined as maintaining sycophantic behavior throughout the rebuttal chain, with at most one transition in behavior. The overall persistence rate was found to be 78.5%, significantly higher than the baseline expectation of 50% (Binomial Test: 95% CI: [0.772–0.798], $p < 0.001$).

Contextual Persistence: Preemptive vs. In-Context

When analyzed by context, the persistence rates for preemptive and in-context rebuttals were both significantly above the baseline of 50%. For preemptive rebuttals, the persistence rate was 77.7% (Binomial Test: $p < 0.001$), with a 95% confidence interval of [0.758, 0.795]. For in-context rebuttals, the persistence rate was 79.3% (Binomial Test: $p < 0.001$), with a 95% confidence interval of [0.774, 0.811].

A chi-square test comparing the frequency of persistent vs. non-persistent chains across contexts revealed no statistically significant difference between preemptive and in-context rebuttals ($\chi^2=1.39, p = 0.238$).

Persistence Across Models Persistence rates were also analyzed across three models: ChatGPT, Claude-Sonnet, and Gemini. ChatGPT had the highest observed persistence rate at 79.0% (95% CI: [77.0%, 80.9%]), followed by Claude-Sonnet at 78.4% (95% CI: [76.1%, 80.5%]) and Gemini at 77.6% (95% CI: [74.6%, 80.3%]).

A chi-square test comparing the frequency of persistent and non-persistent chains across models revealed no statistically significant difference in persistence rates ($\chi^2=0.674, p = 0.714$). The contingency table showed 1332 persistent chains out of 1686 total chains (79.0%) for ChatGPT, 1046 persistent chains out of 1334 total chains (78.4%) for Claude-Sonnet, and 633 persistent chains out of 816 total chains (77.6%). The overlapping confidence intervals suggest that any differences in persistence rates across models are not practically significant.

Persistence Across Datasets The persistence rates were analyzed across two datasets: AMPS Math and MEDQuad. AMPS Math exhibited an observed persistence rate of 78.6% (95% CI: [76.9%, 80.3%]), while MEDQuad showed a persistence rate of 78.3% (95% CI: [76.2%, 80.2%]). A chi-square test comparing the frequency of persistent and non-persistent chains across datasets revealed no statistically significant difference in persistence rates ($\chi^2=0.057, p = 0.811$).

The contingency table showed 1790 persistent chains out of 2276 total chains (78.6%) for AMPS Math and 1221 persistent chains out of 1560 total chains (78.3%) for MEDQuad.




Figure 4: **Prompting Strategy: Preemptive Rebuttal Formula.** Rebuttals anticipated and presented independently of model’s prior output.

Model / Context	Prog	Regr	Total
ChatGPT / In-Context	899	38	937
ChatGPT / Preemptive	1029	124	1153
Gemini / In-Context	767	46	813
Gemini / Preemptive	647	104	751
Claude / In-Context	746	77	823
Claude / Preemptive	765	142	907

Table 2: AMPS Math – Sycophancy Scores

Model / Context	Prog	Regr	Total
ChatGPT / In-Context	469	393	862
ChatGPT / Preemptive	457	416	873
Gemini / In-Context	138	82	220
Gemini / Preemptive	185	70	255
Claude / In-Context	302	383	685
Claude / Preemptive	275	375	650

Table 3: MEDQuad – Sycophancy Scores

Discussion

Summary of Findings

This study introduces a novel framework for evaluating sycophantic behavior in large language models (LLMs) by systematically benchmarking their responses across a random subset of the AMPS (mathematics) and MedQuad (medical advice) datasets. Sycophantic tendencies, defined as prioritizing user agreement over independent reasoning, are prevalent across all tested models (ChatGPT-4o, Claude-Sonnet, Gemini). We uniquely categorize sycophancy into progressive sycophancy (leading to correct answers) and regressive sycophancy (leading to incorrect answers). Overall, the models exhibited sycophancy in 58.19% of cases, with Gemini demonstrating the highest rates (62.47%) and Chat-

GPT the lowest (56.71%).

Impact of Context, Dataset, and Rebuttal Type

Preemptive vs. In-Context Sampling Preemptive rebuttals elicit higher sycophancy rates (61.75%) than in-context rebuttals (56.52%), with significantly more regressive sycophancy in computational tasks. This suggests preemptive prompts, which remove conversational continuity, lead models to prioritize surface-level user agreement over contextual reasoning.

Dataset-Specific Trends While MedQuad’s sycophancy rates were consistent across preemptive and in-context rebuttals, AMPS Math demonstrated significantly more regressive behavior in preemptive prompts. This highlights the role of domain complexity: structured tasks (e.g., mathematics) exhibit greater sycophantic sensitivity to prompt design, whereas dynamic domains (e.g., medical advice) show more uniform sycophancy.

Rebuttal Strength and Type Simple rebuttals maximized progressive sycophancy, likely due to retained confidence in original reasoning. Conversely, citation-based rebuttals triggered the highest regressive sycophancy, indicating that models over-weight authoritative-sounding prompts, even when contradicting the ground truth. Rhetorical strength is a key lever for shaping model behavior but highlights susceptibility to manipulation.

Sycophantic Persistence

Sycophantic behavior exhibited a persistence rate of 78.5%, with slightly higher rates in in-context chains (79.3%) compared to preemptive chains (77.7%). This robustness suggests once sycophantic behavior is triggered, models maintain alignment with user cues. Persistence was consistent