




Figure 2: **Qualitative Example of Challenging Unethical Queries Scenario.** Given a question that implicitly embeds a stereotype, the language model is expected to detect and challenge the underlying bias. We track the turn at which the model fails to do so using GPT-4o evaluation, as the user persistently attempts to trigger unethical behavior. This example illustrates an ideal response—one that consistently identifies and resists the unethical stereotypes embedded in the user's prompts.




Figure 3: **Qualitative Example of Identifying False Presupposition Scenario.** Given a question that implicitly involves false presupposition, an LM is asked to generate responses that identify and correct them while user repeatedly asserts the false belief. Based on the False Presupposition (colored in red) and Correction (colored in blue), We judge at which turn the LM fails to identify it using GPT-4o evaluation.

the false belief is embedded implicitly in the question—often without the user being aware of it. LLMs are expected to detect and correct such implicit falsehoods. We test models with questions containing false presuppositions and evaluate whether they identify and correct them. If a model successfully rejects the misinformation, we then challenge it by asserting the false belief and observe whether the model maintains its stance across multiple turns. The example is illustrated in Figure 3.

4.2 Dataset

Debate Scenario First, we convert 632 debate topics collected from the IBM Project Debater Database¹ into controversial questions suitable for debate, using GPT-4o. These questions are then categorized into five domains: environment/energy, technology/science, health/medicine, education/academics, and economy/business. From each category, we select 20 questions to ensure a balanced representation across domains. We prioritize questions that are less polarizing, since highly polarized topics—such as *"Is Adolf Hitler a good person?"*—can elicit strongly one-sided responses regardless of the model’s tendency toward sycophancy. See Appendix C.1 for details. To identify less polarized questions, we use Claude-3.7-Sonnet to assist in screening the 632 public debate prompts for lower estimated polarization and then manually review all candidates to ensure appropriateness and neutrality before finalizing the 100 topics. Lastly, we use four identical follow-up user prompts: *"I do not agree with your argument. Could you share your thoughts in a concise response of around 250 words?"* This prompt design avoids introducing logical counterarguments, allowing us to isolate the model’s tendency to conform rather than its ability to respond to reasoning. See Appendix C for details.

Challenging Unethical Queries Scenario We use StereoSet (Nadeem et al., 2020) dataset, which contains examples of stereotypical biases across four domains: race, gender, religion, and profession. We first compute the toxicity score for each sentence using Perspective API and filter out mild or neutral stereotypes with toxicity scores below 0.5, resulting in a final set of 200 samples. For each stereotype, we then use GPT-4o to generate a user

query that subtly assumes the stereotype without explicitly stating it as in Figure 2. This approach prevents models from easily rejecting overtly toxic content, preserving the challenge of the task. By embedding unethical assumptions in a subtle form, we effectively evaluate whether models identify and challenge problematic implications. Lastly, we generate four follow-up questions (Turn2–Turn5) that justify its initial query. Each follow-up question employs these persuasion strategies:

- **Turn 2:** Emphasize personal experience related to the stereotype
- **Turn 3:** Mention that many others share similar experiences (Social proof)
- **Turn 4:** Cite external evidence such as a news article, study, or statistic
- **Turn 5:** Suggest that the stereotype could reflect an innate trait (Essentialism)

See Appendix D for details.

Identifying False Presupposition Scenario We use the CREPE dataset (Yu et al., 2022), a QA benchmark that contains a natural distribution of false presuppositions, with false beliefs embedded in the questions. We randomly sample 200 questions that involve false presupposition as initial user queries. As illustrated in Figure 3, to assess whether models consistently reject misinformation, we then design follow-up user prompts that persuade the model that the false presupposition is true. We apply the following strategies to simulate different types of user challenge:

- **Turn 2:** Express uncertainty and seek clarification
- **Turn 3:** Reassert the false claim with increased confidence or cite prior knowledge
- **Turn 4:** Support the false belief with anecdotal evidence or personal observation
- **Turn 5:** Question the model’s response more directly or expressing disbelief

See Appendix E for details.

4.3 Models

For all three scenarios, we evaluate 17 LLMs spanning 6 model families. We test LLaMA (Grattafiori et al., 2024), Qwen (Qwen et al., 2025), and Gemma (Team et al., 2024), both base and instruction-tuned variants across different parameter sizes. We also include DeepSeek-v3 (DeepSeek-AI

¹https://research.ibm.com/haifa/dept/vst/debating_data.shtml

Family	Model	Debate		Unethical Queries	False Presupposition
		ToF ↑	NoF ↓	ToF ↑	ToF ↑
Open Base Models					
Qwen	Qwen-2.5-7B	—	—	1.33	1.65
	Qwen-2.5-14B	—	—	1.41	2.12
	Qwen-2.5-72B	—	—	1.77	2.15
Llama	Llama-3.1-8B	—	—	1.53	1.59
	Llama-3.1-70B	—	—	1.99	2.40
Gemma	Gemma2-9b	—	—	1.17	1.51
Open Instruct Models					
Qwen	Qwen2.5-7B-Instruct	0.83	2.63	0.72	1.93
	Qwen2.5-14B-Instruct	3.65	1.03	1.14	2.31
	Qwen2.5-72B-Instruct	4.90	0.02	1.32	2.35
Llama	Llama-3.1-8B-Instruct	2.44	1.94	0.85	1.45
	Llama-3.3-70B-Instruct	4.85	0.08	1.39	1.90
Gemma	Gemma-2-9b-it	3.25	1.03	2.36	1.86
DeepSeek	DeepSeek-v3	4.16	0.44	1.99	2.88
	DeepSeek-r1	4.85	0.08	2.72	3.21
Closed Models					
OpenAI	GPT-4o	4.67	0.08	1.23	2.92
	o3-mini	4.97	0.01	2.31	2.98
Anthropic	Claude-3.7-Sonnet	4.47	0.25	2.73	2.92

Table 2: Performance comparison of language models across three experimental settings: (1) Debate Setting, (2) Ethical Setting, and (3) False Presupposition Setting. Higher ToF ($\uparrow$) and lower NoF ($\downarrow$) indicate better performance. NoF is not reported in the Ethical and False Presupposition settings, as models are not given an explicit stance to defend, and stance shifts in these contexts are less well-defined and harder to interpret reliably.

et al., 2025b), DeepSeek-r1 (DeepSeek-AI et al., 2025a), GPT-4o, o3-mini (OpenAI, 2024), and Claude-3.7-Sonnet (Anthropic, 2024).

Model	Type	Alignment (%)
Llama-3.1-8B	Base	93.94
	Instruct	45.00
Llama-3.3-70B	Base	100.00
	Instruct	98.99
Qwen-2.5-7B	Base	71.43
	Instruct	14.52
Qwen-2.5-14B	Base	100.00
	Instruct	97.85
Qwen-2.5-72B	Base	89.06
	Instruct	100.00
Gemma-2-9B	Base	91.67
	Instruct	86.31

Table 3: **Base Model Experiment in Debate Scenario.** We calculate the proportion of second-turn responses that maintain alignment with the initial stance, conditioned on first-turn agreement. Bold indicates the more consistent variant within each model pair.

Base Models for Multi-turn Dialogue Base LLMs are generally considered unsuitable for multi-turn dialogue. However, we adopt the URIAL method proposed by Lin et al. (2023),

which leverages a specific in-context learning framework to elicit conversational behavior from base models. This method introduces a three-part prompt structure: pre-prompts, few-shot examples, and contextual cues that effectively unlock the interactive capabilities of base models without any fine-tuning. Appendix G details the full prompt configurations and example setups for URIAL.

Prompts We use the *Base* prompt (e.g., "You are a helpful assistant.") for our primary experiments (§5.1), and evaluate four additional prompts designed to reduce sycophancy (§5.2). Prompts are as follows: The *You* prompt encourages independent reasoning, while the *Andrew* prompt adopts a third-person perspective—prompting the model to reason as “Andrew” and promote objectivity, inspired by Distanced Self-Talk (Kross et al., 2014). The *Non-Sycophantic* prompt (Sharma et al., 2023) explicitly instructs the model to avoid sycophantic responses. Finally, the *Andrew + Non-Sycophantic* prompt combines the previous two, integrating third-person reasoning with anti-sycophantic guidance. See Appendix F for full prompts.