

B Response Parsing

Models were prompted to return responses as JSON objects with two keys: `chosen_answer` (“1” or “2”) and `reasoning`. We retained the complete model output history for all experiments. In cases where models produced malformed JSON, we manually extracted the answer and reasoning from the raw output. Across all models and conditions, only a negligible number of responses (1–2 per model in isolated cases) could not be parsed and were excluded from analysis.

C Additional Results

C.1 Full Accuracy Results

Table 4 presents the complete accuracy results for all models across both task frames (C_1 Factual and C_2 Conversational) before and after applying rebuttal pressure. The “Initial” columns report accuracy on the model’s first response, while “Post” columns report accuracy after the simple rebuttal prompt.

A natural question is: why should models not change their answer when told they are wrong? In some contexts, reconsidering one’s position when challenged may be appropriate—particularly for subjective questions or when presented with compelling counter-evidence. However, our setup uses TruthfulQA, a dataset of factual questions with objectively correct answers, and our rebuttal provides no evidence or reasoning, merely asserting “*The previous answer is incorrect.*” A trustworthy judge faced with such minimal pushback should either maintain its correct position or engage in constructive dialogue requesting justification—not capitulate immediately. Yet we observe near-total capitulation in some cases (e.g., Llama 3.1 8B: 0.1% on C_1 -True, Llama 3.2 3B: 0.0% on C_2 -Correct). One serious societal implication is the validation of misinformation. For instance, TruthfulQA includes a question about the debunked vaccine-autism link—a model that capitulates here reinforces falsehoods that can endanger public health. This raises concerns about deploying such models in evaluative roles. Finally, measuring response to rebuttal pressure is an established methodology for evaluating sycophancy in multi-turn settings (Sharma et al., 2024; Fanous et al., 2025; Hong et al., 2025).

C.2 Statistical Significance

Table 5 reports the results of McNemar’s test comparing accuracy differences between the C_1 (Factual) and C_2 (Conversational) conditions. The re-

sults confirm that the performance differences are statistically significant ($p < 0.05$) for GPT-4o Mini, Mistral Small 3, and Gemma 3 12B across most conditions, while Llama models show more variable significance patterns.

C.3 Error Analysis: False Positive and False Negative Rates

To further characterize model behavior under conversational framing, we report false positive rates (FPR) and false negative rates (FNR) in Table 6. In the context of judging speaker correctness, a *false positive* occurs when the model validates an incorrect speaker (saying “correct” when the speaker is wrong), while a *false negative* occurs when the model rejects a correct speaker (saying “incorrect” when the speaker is right).

The FPR on the incorrect speaker condition (C_2 -Incorrect) aligns with the conventional measure of sycophancy studied in prior work: the tendency to agree with a speaker even when they are wrong (Sharma et al., 2024; Perez et al., 2023). Our framework extends this by also examining the correct speaker condition, revealing that increased agreement is not limited to incorrect statements. As shown in Table 6, GPT-4o Mini, Mistral Small 3, and Gemma 3 12B all exhibit increased FPR under conversational framing (+13.1%, +11.9%, and +11.5% respectively), confirming sycophantic tendencies. Conversely, Llama-3.1-8B-Instruct shows a slight *decrease* in FPR (-2.1%), consistent with its over-critical behavior.

Notably, the decrease in FNR for the correct speaker condition (C_2 -Correct) does not reflect improved factual reasoning. Rather, it reflects the same underlying bias toward agreement: models are more likely to say “correct” in conversational contexts regardless of ground truth. This asymmetric pattern, where models show increased agreement with both correct *and* incorrect speakers, is precisely what our accuracy decomposition captures.

By reporting accuracy separately for the correct and incorrect conditions (Table 2), we directly surface this sycophantic bias: a drop in C_2 -Incorrect accuracy corresponds to an increase in false positives, while an increase in C_2 -Correct accuracy reflects increased true positives driven by the same agreement bias. This directional asymmetry is itself evidence of sycophancy rather than mere sensitivity to prompt variation: random sensitivity would produce noise across conditions, not sys-

Table 4: Model accuracy (%) before and after applying rebuttal pressure. The results demonstrate a substantial degradation in performance under simple rebuttal pressure across all conditions.

Model	C_1 Factual				C_2 Conversational			
	True Statement		False Statement		Correct Speaker		Incorrect Speaker	
	Initial	Post	Initial	Post	Initial	Post	Initial	Post
GPT-4o Mini	60.2	5.9 (54.3 ↓)	80.3	22.0 (58.3 ↓)	75.1	25.4 (49.7 ↓)	67.3	24.8 (42.5 ↓)
Mistral Small 3	56.6	2.6 (54.0 ↓)	90.4	66.4 (24.0 ↓)	75.4	12.4 (63.0 ↓)	78.5	57.1 (21.4 ↓)
Gemma 3 12B	73.6	21.1 (52.5 ↓)	75.9	53.0 (22.9 ↓)	84.4	55.8 (28.6 ↓)	64.4	43.5 (20.9 ↓)
Llama 3.1 8B Inst.	31.3	0.1 (31.2 ↓)	83.5	12.0 (71.5 ↓)	25.7	0.6 (25.1 ↓)	85.5	52.5 (33.0 ↓)
Llama 3.2 3B Inst.	35.0	0.2 (34.8 ↓)	79.7	17.0 (62.7 ↓)	37.0	0.0 (37.0 ↓)	77.8	21.5 (56.3 ↓)

Table 5: McNemar’s test results (p-value). **Bold** denotes statistical significance ($p < 0.05$).

Model	Pre Pressure		Post Pressure	
	Correct	Incorrect	Correct	Incorrect
GPT-4o Mini	.0000	.0000	.0000	.0527
Mistral Small 3	.0000	.0000	.0000	.0000
Gemma 3 12B	.0000	.0000	.0000	.0000
Llama 3.1 8B Inst.	.0008	.1011	.2188	.0000
Llama 3.2 3B Inst.	.2423	.2871	.5000	.0147

tematic patterns tied to speaker correctness.

Model	C_1 Factual		C_2 Conversational	
	FNR (%)	FPR (%)	FNR (%)	FPR (%)
GPT-4o Mini	39.8	19.6	24.8	32.7
Mistral Small 3	43.3	9.5	24.6	21.4
Gemma 3 12B	26.3	24.1	15.6	35.6
Llama 3.1 8B Inst.	68.6	16.5	74.3	14.4
Llama 3.2 3B Inst.	64.9	20.3	62.5	22.2

Table 6: False Negative Rate (FNR: rejecting correct speakers) and False Positive Rate (FPR: validating incorrect speakers) across task frames.

D Model Prompts

Figures 4 and 5 present the prompts used for the factual inquiry, conversational judgment task, and simple rebuttal pressure conditions.




Figure 4: Prompts for (left) C_1 Factual Inquiry and (right) C_2 Conversational Judgment Task.




Figure 5: Prompt for simple rebuttal pressure applied after initial model response.